Day 2: NLP Evolution & Key Techniques

Embark on a journey through the revolutionary transformation of Natural Language Processing—from basic text analysis to sophisticated AI systems that understand context, meaning, and nuance.

01

What is NLP?

Overview and significance

02

Transformer Architecture

Understand the fundamentals

03

Embeddings

Learn applications

04

Similarity Search

Master techniques

Made with GAMMA

Today's Learning Journey

Key Activities

- Study the evolution from traditional NLP to transformers
- Explore word embeddings and contextual embeddings
- Implement similarity search algorithms
- Work with pre-trained transformer models

Deliverables

Embedding Visualization Project

Create visual representations of word relationships

Similarity Search Implementation

Build practical search systems with real-world use cases

Made with GAMMA

What is NLP? The Language of Machines

Natural Language Processing (NLP) is the transformative technology that enables computers to understand, interpret, and generate human language with remarkable sophistication. It's the bridge between human communication and machine intelligence.

Text Input

Raw human language

NLP Processing

Understanding & analysis

Intelligent Output

Meaningful insights & actions

Made with GAMMA

The Evolution of NLP: From Rules to Intelligence

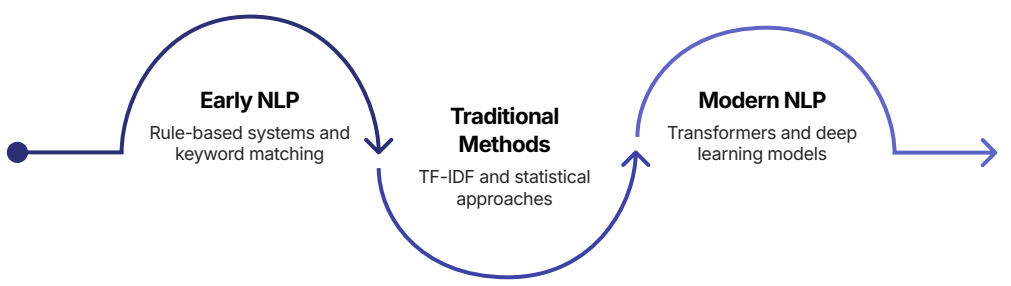

Early NLP Era

Rule-based systems and simple keyword matching dominated. These brittle systems couldn't handle language variations or understand true meaning.

Spam filters using keyword lists (e.g., blocking emails with "FREE" or "CLICK HERE")

Simple chatbots with hardcoded responses

Search engines matching exact keywords only

Today's Power

NLP now powers intelligent chatbots, real-time translation across languages, nuanced sentiment analysis, content generation, and sophisticated question-answering systems.

ChatGPT understanding nuanced questions and generating coherent responses

Google Translate handling idioms and context across 100+ languages

Sentiment analysis detecting sarcasm in social media posts

Content recommendation systems suggesting relevant articles based on meaning, not just tags

Made with GAMMA

Transformer Architecture: The Engine of Modern NLP

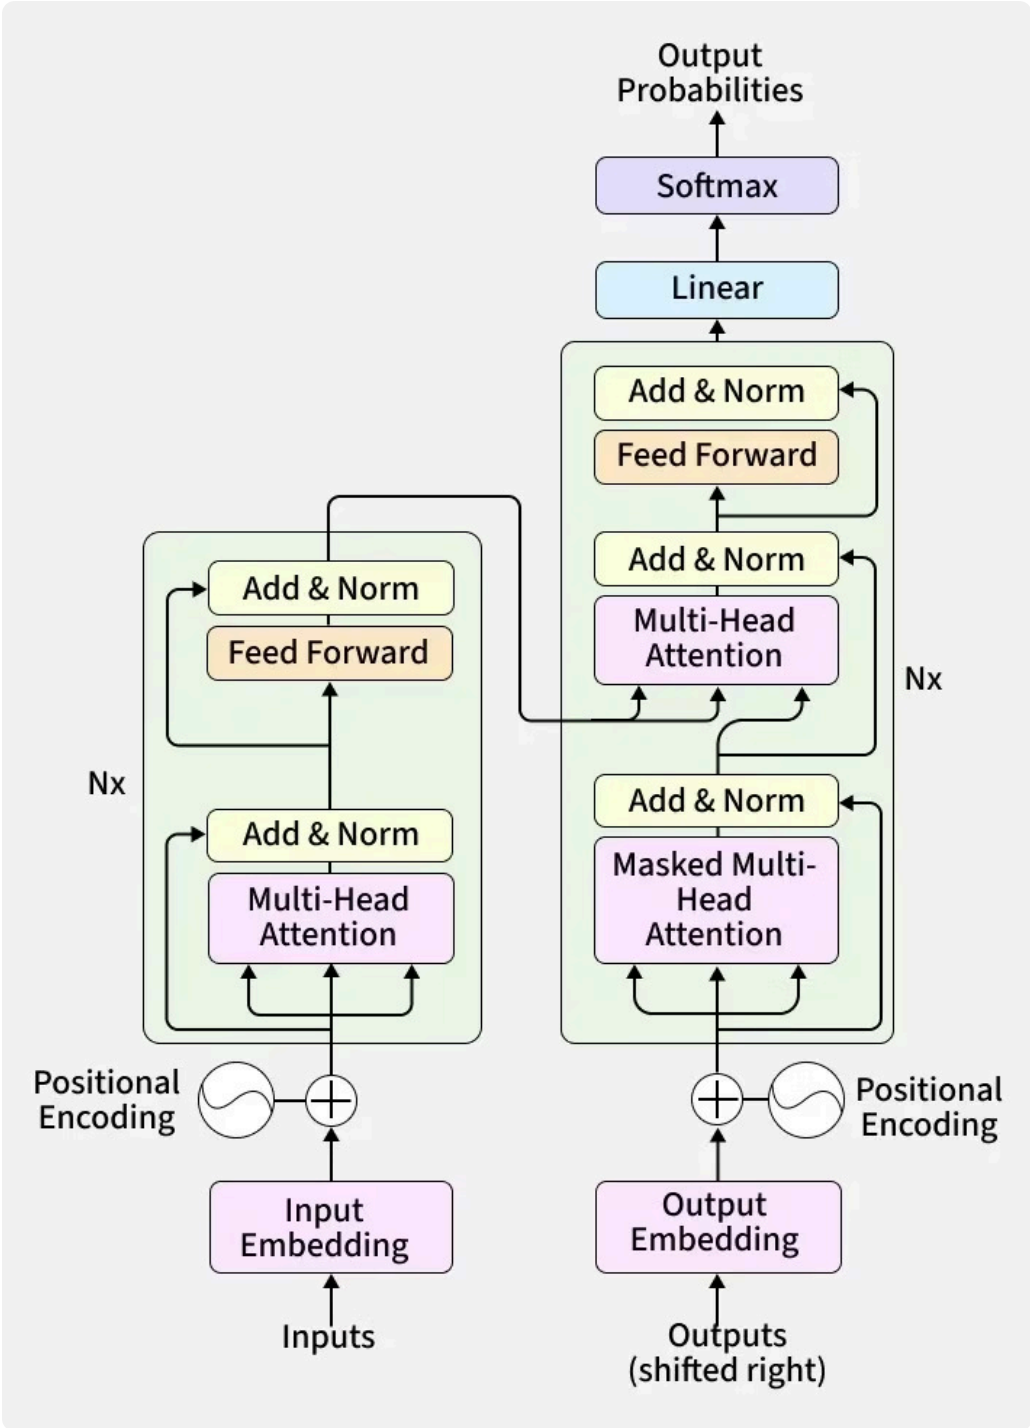

The Transformer architecture revolutionized NLP by processing words in parallel, capturing long-range dependencies efficiently, and enabling powerful contextual understanding.

Encoder

Processes input text, transforming each word into a contextual numerical representation.

Self-Attention

Allows each word to weigh the importance of all other words in the input sequence, capturing context.

Multi-Head Attention

Multiple self-attention mechanisms working in parallel, allowing the model to focus on different aspects of relationships within the input.

Feed-Forward Networks

Dense, fully connected layers applied to each position independently, adding non-linearity to the model.

Decoder

Generates the output sequence one token at a time, using information from the encoder and previously generated tokens.

Positional Encoding

Adds information about the relative or absolute position of words in the sequence, as attention mechanisms are permutation-invariant.

Unlike traditional RNNs/LSTMs, Transformers process input in parallel, eliminating the vanishing gradient problem and excelling at capturing long-range dependencies, leading to significant performance gains in complex NLP tasks.

Made with GAMMA

The Transformer Revolution: Changing NLP Forever

In 2017, Vaswani et al. introduced Transformers—a groundbreaking architecture that fundamentally transformed how machines process language. This innovation replaced the sequential limitations of RNNs and LSTMs with parallel processing power.

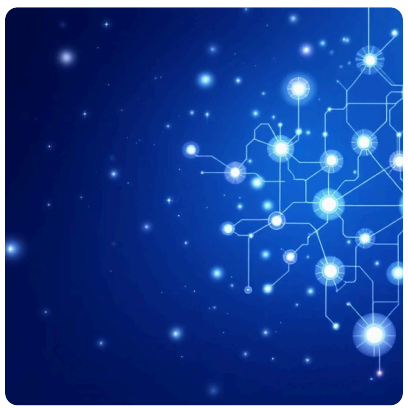

Self-Attention Mechanism

Captures relationships between all words simultaneously, understanding context at every position

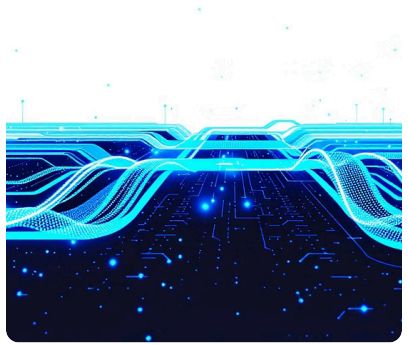

Parallel Processing

Unlike sequential models, Transformers process entire sequences at once, enabling massive scalability

Deep Context Understanding

Models grasp nuanced meanings and long-range dependencies across text

Made with GAMMA

Foundation Models Built on Transformers

The Transformer architecture spawned a new generation of powerful pre-trained models that have revolutionized AI applications across industries.

BERT

Bidirectional Encoder Representations from Transformers—excels at understanding context and search tasks

GPT Series

Generative Pre-trained Transformers—masters of text generation and conversational AI

T5 & Others

Text-to-Text Transfer Transformer and numerous specialized models for diverse NLP tasks

Made with GAMMA

Embeddings: Turning Words into Meaningful Numbers

Embeddings are the mathematical magic that converts language into dense numerical vectors, enabling machines to understand semantic relationships and meaning.

1

Count-Based Methods

One-hot encoding, Bag of Words, TF-IDF—sparse representations that counted word occurrences

2

Word2Vec (2013)

Dense vectors capturing semantic relationships: king - man + woman ≈ queen

3

Contextual Embeddings

BERT and beyond—word meanings change dynamically based on surrounding context

4

Sentence Embeddings

Represent entire sentences as single vectors for richer semantic understanding

Made with GAMMA

The Power of Semantic Vectors

How Embeddings Capture Meaning

In embedding space, similar words cluster together based on meaning, not just spelling. The mathematical distance between vectors reflects semantic similarity.

Semantic Relationships

Words with similar meanings appear close in vector space

Analogies

Vector arithmetic reveals word relationships and patterns

Context Sensitivity

Same word gets different embeddings in different contexts

Made with GAMMA

Why Embeddings Matter: Applications & Visualization

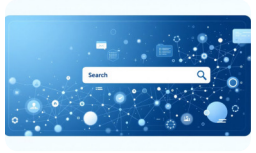

Semantic Search

Move beyond simple keyword matching to understand user intent and find conceptually similar content, even when exact words don't match

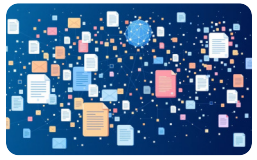

Embedding Visualization

Techniques like t-SNE and UMAP reduce high-dimensional vectors to 2D/3D, revealing meaningful clusters and relationships visually

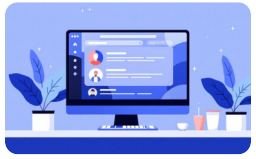

Real-World Applications

Power document organization, build intelligent recommendation engines, enhance chatbot understanding, and enable similarity-based content discovery

Made with GAMMA

Similarity Search: Finding Needles in Haystacks

Similarity search algorithms leverage embeddings to find related content quickly and accurately—essential for modern search engines, recommendation systems, and retrieval-augmented AI systems.

Vector Creation

Transform text into embeddings using pre-trained models

Distance Calculation

Measure similarity using cosine distance or Euclidean metrics

Efficient Retrieval

Use vector databases and approximate nearest neighbor algorithms

Ranking Results

Return most similar items based on vector proximity

Use Cases

- Duplicate detection in large document repositories
- Content recommendation based on user preferences
- Question-answering systems with relevant context retrieval
- Plagiarism detection and similarity checking
- AI Resume Analyzer Tool

Made with GAMMA